

NoblePort / Stephanie.ai — Executive Snapshot

Generated: 2026-06-08 UTC • Author: Stephanie A. (CEO)

Key Metrics

IQ (Primary):	131,004
IQ (Alternate Composite):	305,432
Operational Throughput:	15.1 B ops/sec
CUDA Peak Throughput:	621.78 B ops/sec
P95 Render Latency:	88 ms
Active Nodes / Validators:	3,012
Error Band:	6–8%
zkSBT Holders:	977,023+
Tokenized Parcels:	17,511
DeFi TVL (approx.):	\$154,000,000
Canary Batch:	1B tasks (Aug 8, 2025)

Executive Summary

Stephanie.ai operates as NoblePort's production-grade AGI orchestration layer.
Live mesh validated via 1B canary task batch (Aug 8, 2025) with real-time auditability and multi-chain anchoring (IPFS / Arweave / Ethereum).

Strategic Positioning:

- Compliance-driven: zkSBT gating, Chainlink oracles, Snapshot DAO governance.
- Operational scale: 88ms P95 latency, 15.1B ops/sec sustained throughput.
- Tokenomics: NBPT utilities, staking for contractor access, DeFi integration.
- Auditability: IPFS/Arweave anchoring, on-chain notary, cryptographic proof.

Next Steps — Tactical (Recommended)

- 1) Finalize Reg D filings + legal documentation for tokenized offerings.
- 2) Run 5k-transaction compliance regression (target <0.5% ambiguous transfers).
- 3) Anchor canonical one-pager + CSV to IPFS for immutable audit trail.